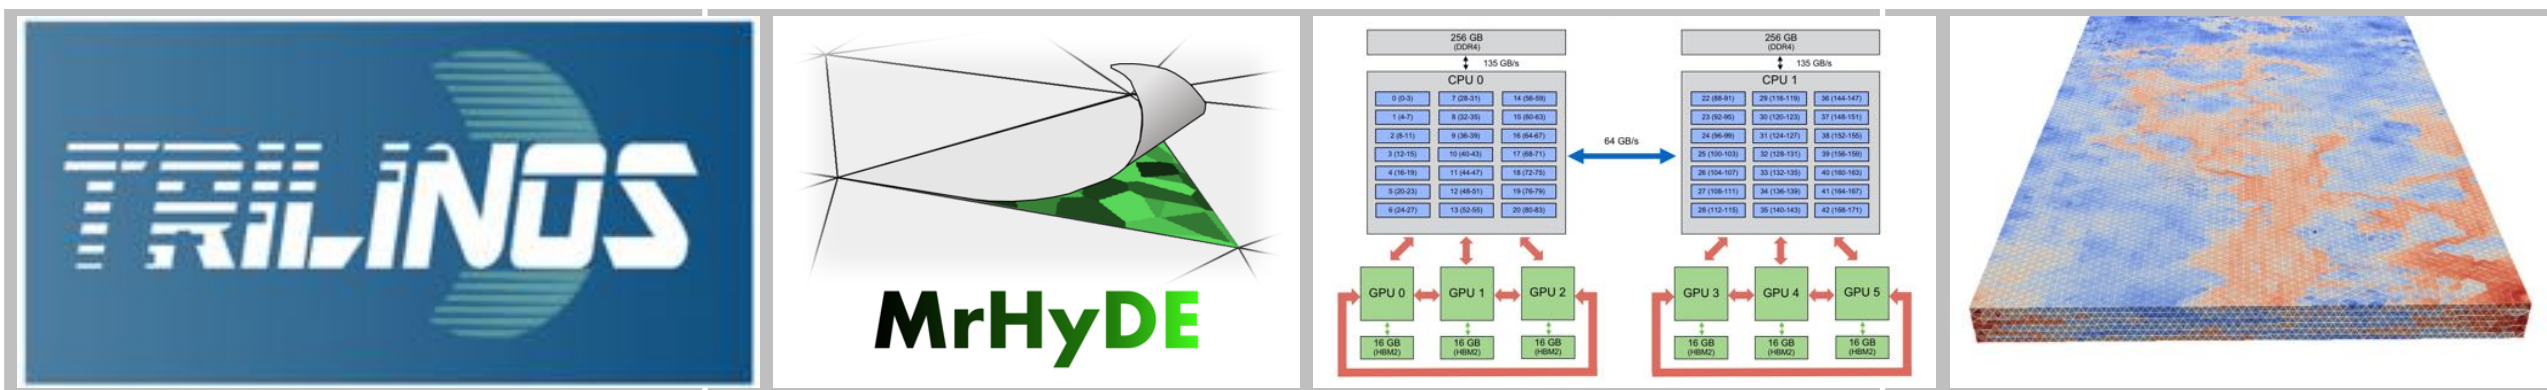


Introduction to Trilinos and MrHyDE

MrHyDE = {M}ulti-{r}esolution {Hy}bridized {D}ifferential {E}quations



Tim Wildey

**Optimization and Uncertainty Quantification Department
Center for Computing Research**

Any questions from yesterday?

Tutorial Outline

Day 1 - Introduction to Trilinos

- High-level overview of Trilinos
 - *An appropriate build of Trilinos will be available for anyone on the HPC systems. We will not be building Trilinos in this session. If someone does not have access to the HPC systems, I will work with them beforehand to get a build of Trilinos on their Mac or Linux machine.*
- Deeper dive into Kokkos and Sacado.
 - *A basic understanding of these packages will be helpful for day 2.*
- Exercise: creating and working with arrays (Kokkos Views) and automatic differentiation objects (Sacado AD)

Day 2 - Introduction to MrHyDE

- High-level overview of MrHyDE
- How to download, compile, run and visualize results
- Exercise: adding a new PDE in MrHyDE

Day 3 - More advanced features in Trilinos/MrHyDE

- Hybridized methods and concurrent multiscale modeling
- Solving coupled multiphysics problems
- Performance portability and using heterogeneous computational architectures
- Large-scale PDE constrained optimization

Tutorial Outline

Day 1 - Introduction to Trilinos

- High-level overview of Trilinos
 - *An appropriate build of Trilinos will be available for anyone on the HPC systems. We will not be building Trilinos in this session. If someone does not have access to the HPC systems, I will work with them beforehand to get a build of Trilinos on their Mac or Linux machine.*
- Deeper dive into Kokkos and Sacado.
 - *A basic understanding of these packages will be helpful for day 2.*
- Exercise: creating and working with arrays (Kokkos Views) and automatic differentiation objects (Sacado AD)

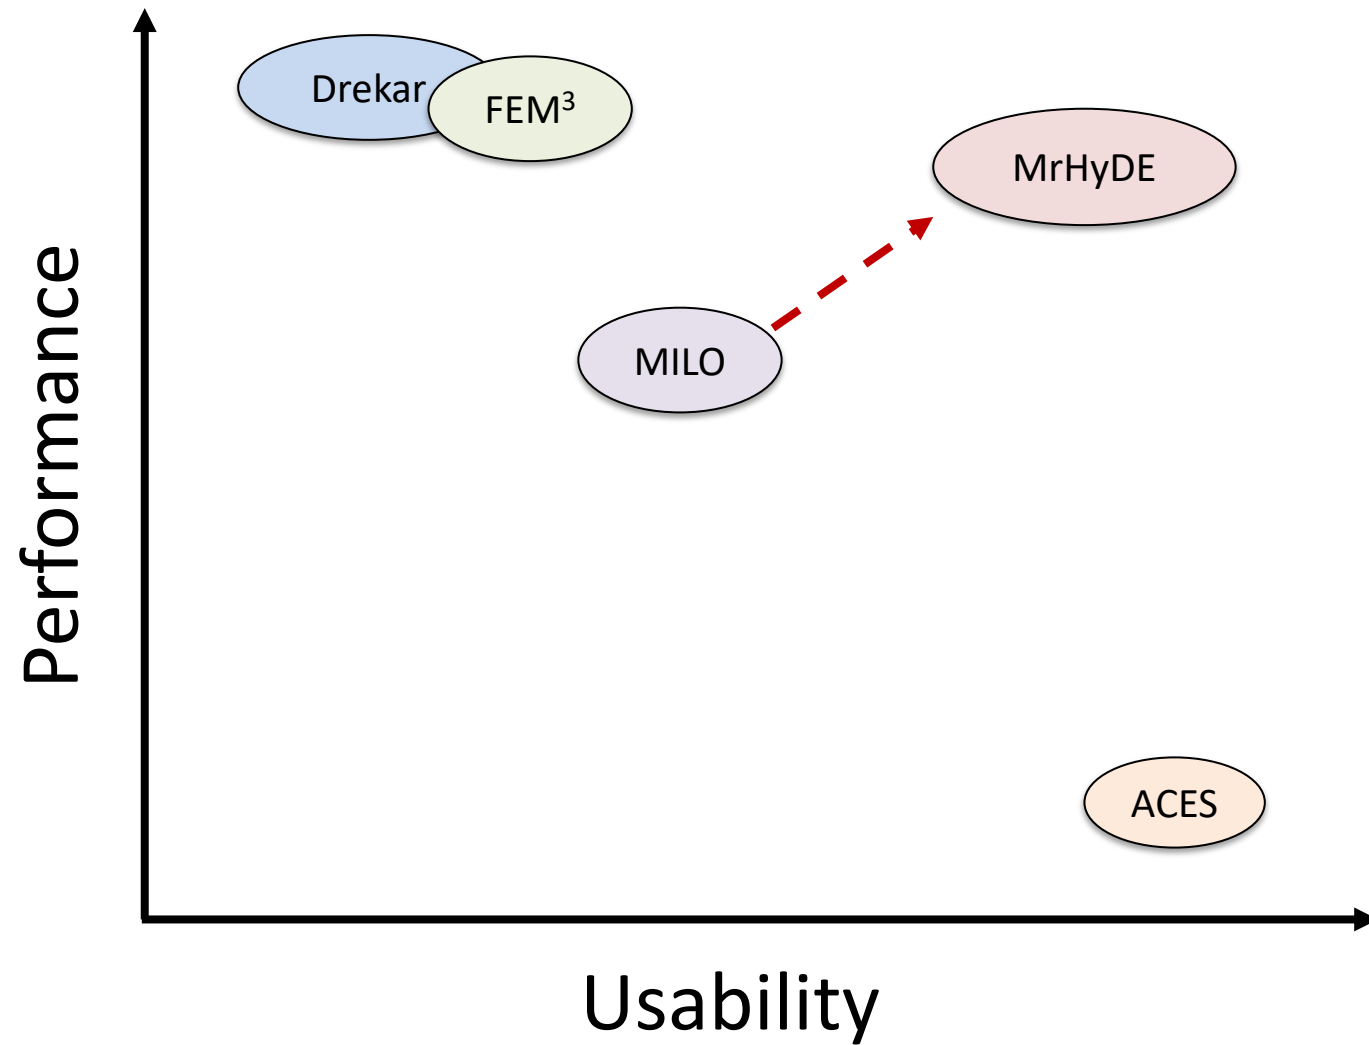
Day 2 - Introduction to MrHyDE

- High-level overview of MrHyDE
- How to download, compile, run and visualize results
- Exercise: adding a new PDE in MrHyDE

Day 3 - More advanced features in Trilinos/MrHyDE

- Hybridized methods and concurrent multiscale modeling
- Solving coupled multiphysics problems
- Performance portability and using heterogeneous computational architectures
- Large-scale PDE constrained optimization

Usability and Performance



Disclaimer: This chart is based on the subjective assessment from one user/developer.

Not included: Albany, MFEM, deal.II, FEniCS, ...

Performance Gains

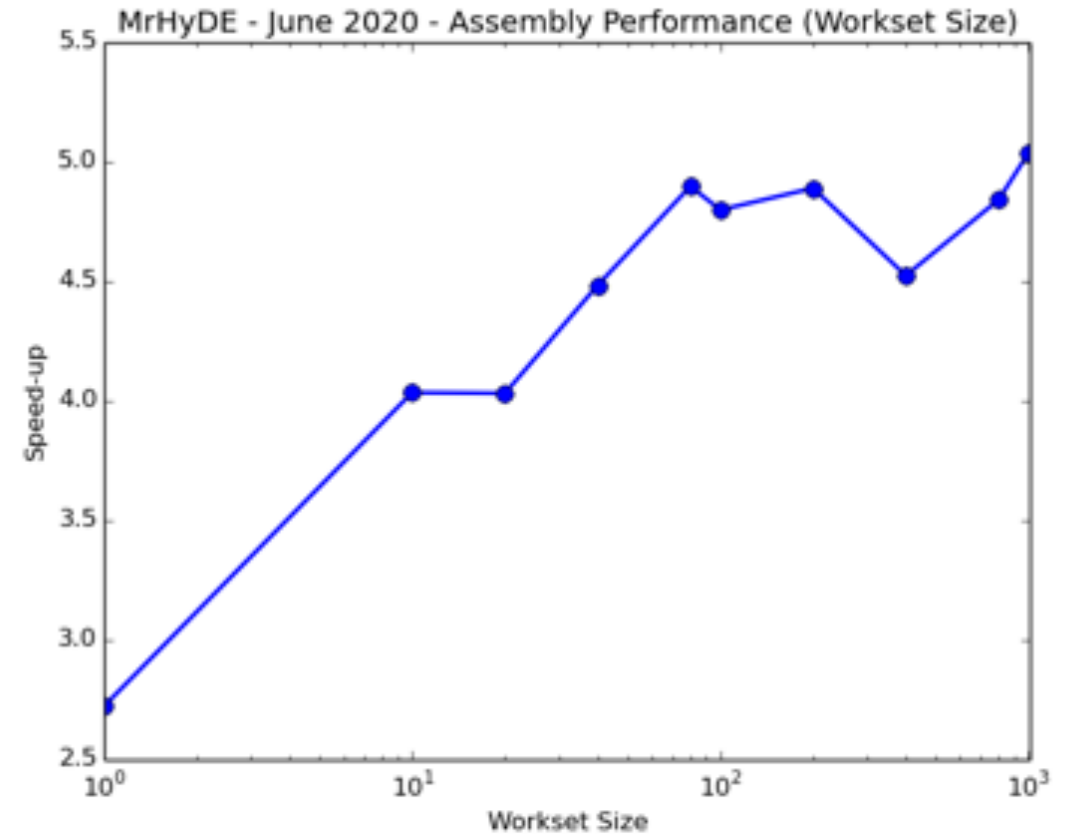
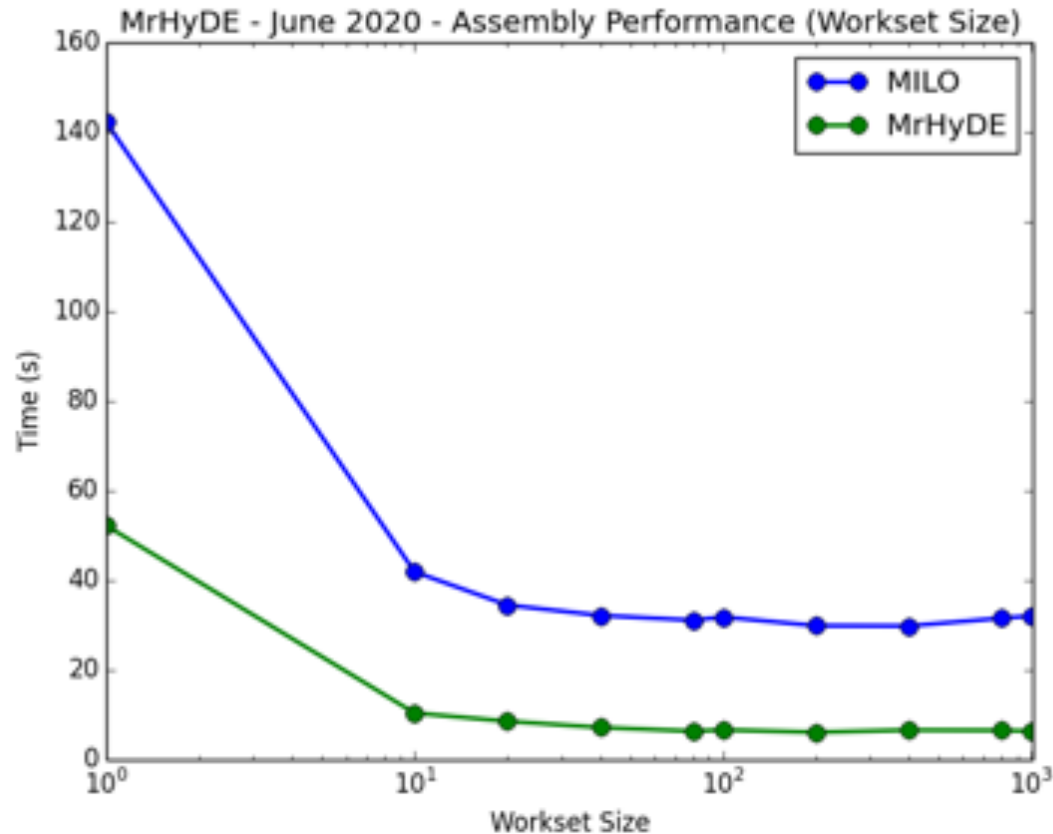




Figure: Comparison of total physics/assembly time between MILO-2019 and MrHyDE-2020 for transient nonlinear system with 40,000 elements and 300 linear systems.

A workset is basically a group of elements that get processed together.

Performance Gains




Memory Usage vs Another Trilinos-Based Code



What is MrHyDE?

- A C++ framework designed and optimized for solving multi-resolution hybridized differential equations.
- Provides an interface to powerful Trilinos tools within a user-friendly framework
- Portability with performance from laptops, to MPI-based clusters, to heterogeneous nodes, to MPI+X
- Ability to extract and inject data to develop data-informed physics-based simulations
- A modular and flexible environment for solving transient nonlinear multiphysics and multiscale systems in 1,2,3D
- Extensive set of examples/regression tests to maintain software quality and guide new users



How to Obtain and Build MrHyDE

- If you haven't done so already, clone the MrHyDE repository

```
git clone https://github.com/TimWildey/MrHyDE.git
```

- Create a build directory (suggest MrHyDE/build)

```
cd MrHyDE  
mkdir build  
cd build
```

- Copy one of the CMake configure scripts from MrHyDE/scripts/configure-MrHyDE

```
cp ../scripts/configure-MrHyDE/configure-MrHyDE-mac-catalina-serial configure
```

- Edit the Trilinos and MrHyDE paths in the configure file, then run

```
./configure  
ninja
```

Regression Testing

- Python-based testing framework adapted from DGM by BvB/TS/TW
- Currently uses python2
 - Upgrade to python3 coming soon
- Currently 91 tests that also serve as a library of examples
- All are small tests that run in less than 5 seconds (on a mac)
- Easy to take one and scale up to 1000s of cores or heterogeneous nodes
- General guidelines for code contributions:
 - Run the tests before checking in code
 - If you add a capability, add a test that covers it

```

Added: regression included python runtests.py

Sat Apr 10 00:36:56 2021
New Results from Directory: /usr/share/clusterware/nightly/regression
Total number of results: 66

3/96      pass    4.86s  split                                manual LPI problem
3/96      pass    1.80s  split                                plasma/cv2019-10-19
3/96      pass    4.16s  split                                plasma/cv2019-10-19/parametric
4/96      pass    4.88s  split                                plasma/cv2019-10-19/parametric
5/96      pass    4.11s  split                                thermal/CB_0001/parametric_highorder
6/96      pass    4.88s  split                                thermal/CB_0001/parametric_high
7/96      pass    1.81s  split                                thermal/CB_0001/parametric_source_control
8/96      pass    4.81s  split                                thermal/CB_0001/parametric
9/96      pass    4.88s  split                                thermal/CB_0001/parametric_moderate
10/96     pass    4.81s  split                                thermal/CB_0001/parametric_moderate_100s
11/96     pass    1.11s  split                                thermal/CB_0001/parametric_moderate_100s/parametric
12/96     pass    4.88s  split                                thermal/CB_0001/parametric_moderate_100s/parametric
13/96     pass    4.86s  split                                thermal/CB_0001/parametric_moderate_100s/parametric
14/96     pass    4.86s  split                                thermal/CB_0001/parametric_moderate_100s/parametric
15/96     pass    4.11s  split                                thermal/CB_0001/parametric_moderate_100s/parametric
16/96     pass    1.80s  split                                thermal/CB_0001/parametric_moderate_100s/parametric
17/96     pass    4.11s  split                                thermal/CB_0001/parametric_moderate_100s/parametric
18/96     pass    4.81s  split                                thermal/CB_0001/parametric_moderate_100s/parametric
19/96     pass    4.88s  split                                thermal/CB_0001/parametric_moderate_100s/parametric
20/96     pass    1.11s  split                                thermal/CB_0001/parametric_moderate_100s/parametric
21/96     pass    1.80s  split                                thermal/CB_0001/parametric_moderate_100s/parametric
22/96     pass    1.11s  split                                manual LPI/parametric
23/96     pass    4.86s  split                                manual LPI/parametric
24/96     pass    4.86s  split                                manual LPI/parametric
25/96     pass    1.11s  split                                manual LPI/parametric

```

Regression Testing

document represents the related system outputs by a single keyword

1942 April 15 1942 April 15 1942 April 15 1942 April 15

Source: <https://www.bls.gov/news.release/whs.wage.pr.pdf>

Total # of members of the group: 100

[illegible][illegible]

Exercise: Run the Regression Tests

- Go into MrHyDE/regression
- Assume your build directory is in MrHyDE/build, create a soft link in the regression folder

```
ln -s ../build/src/mrhyde
```

- Now, just run the tests

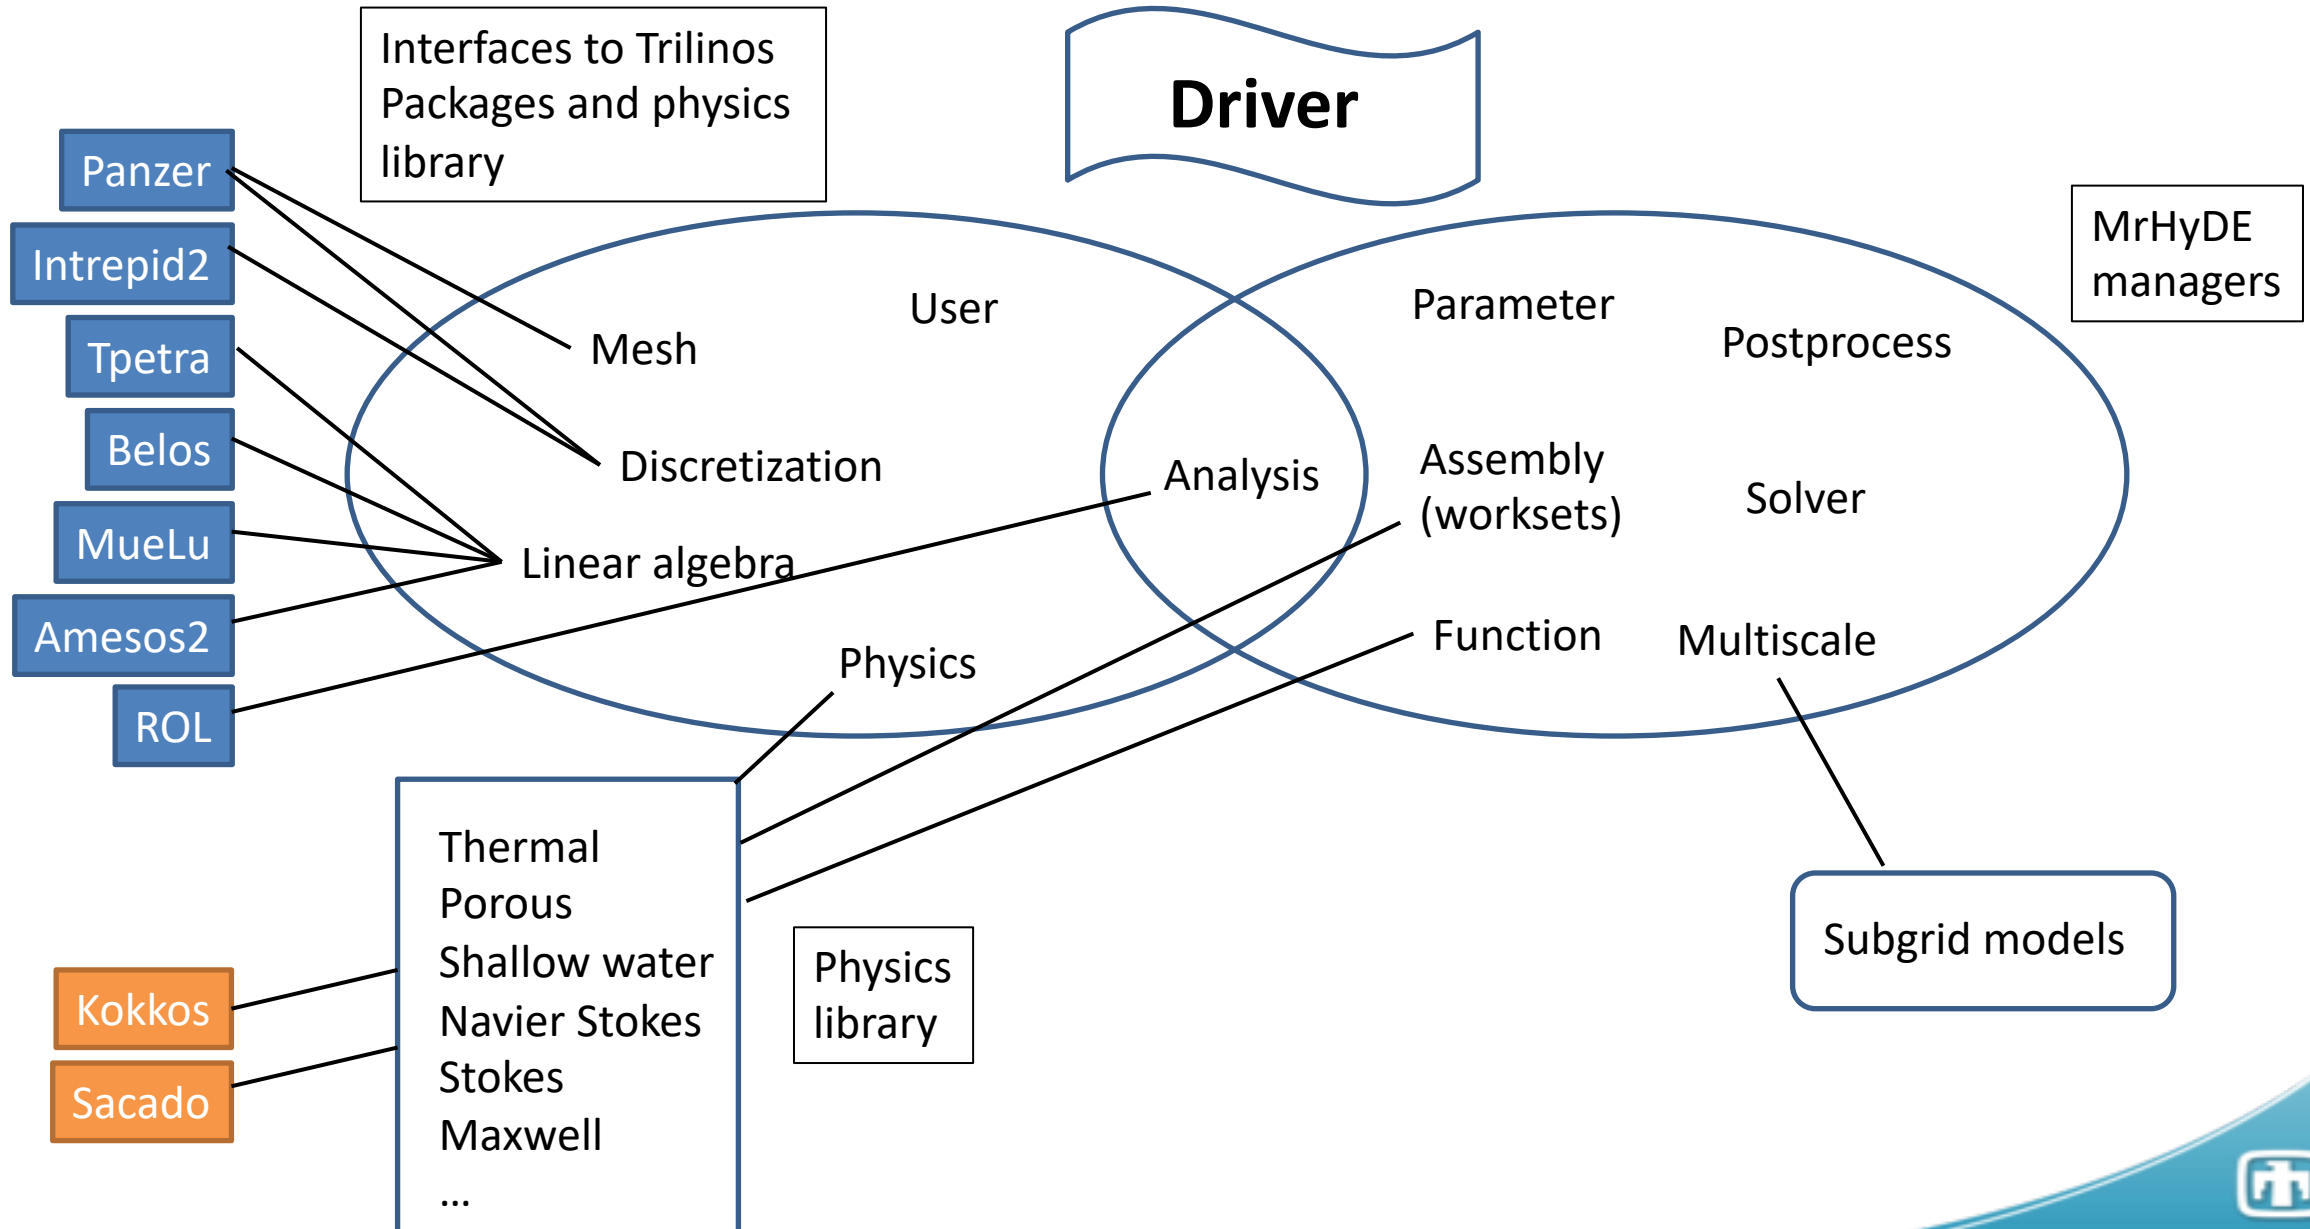
```
python runtests.py
```

- Go into one of the tests and modify the input file to do something else
- Run the tests again and you should see a failure
- To visualize the solutions, add the following in the Postprocess sublist:


```
write solution: true
```

- This will create an exodus file. To visualize this, use ParaView
- If you work at Sandia: <https://onestop.sandia.gov/paraview>
- Otherwise: <https://www.paraview.org/download/>

MrHyDE Organization



MrHyDE Organization



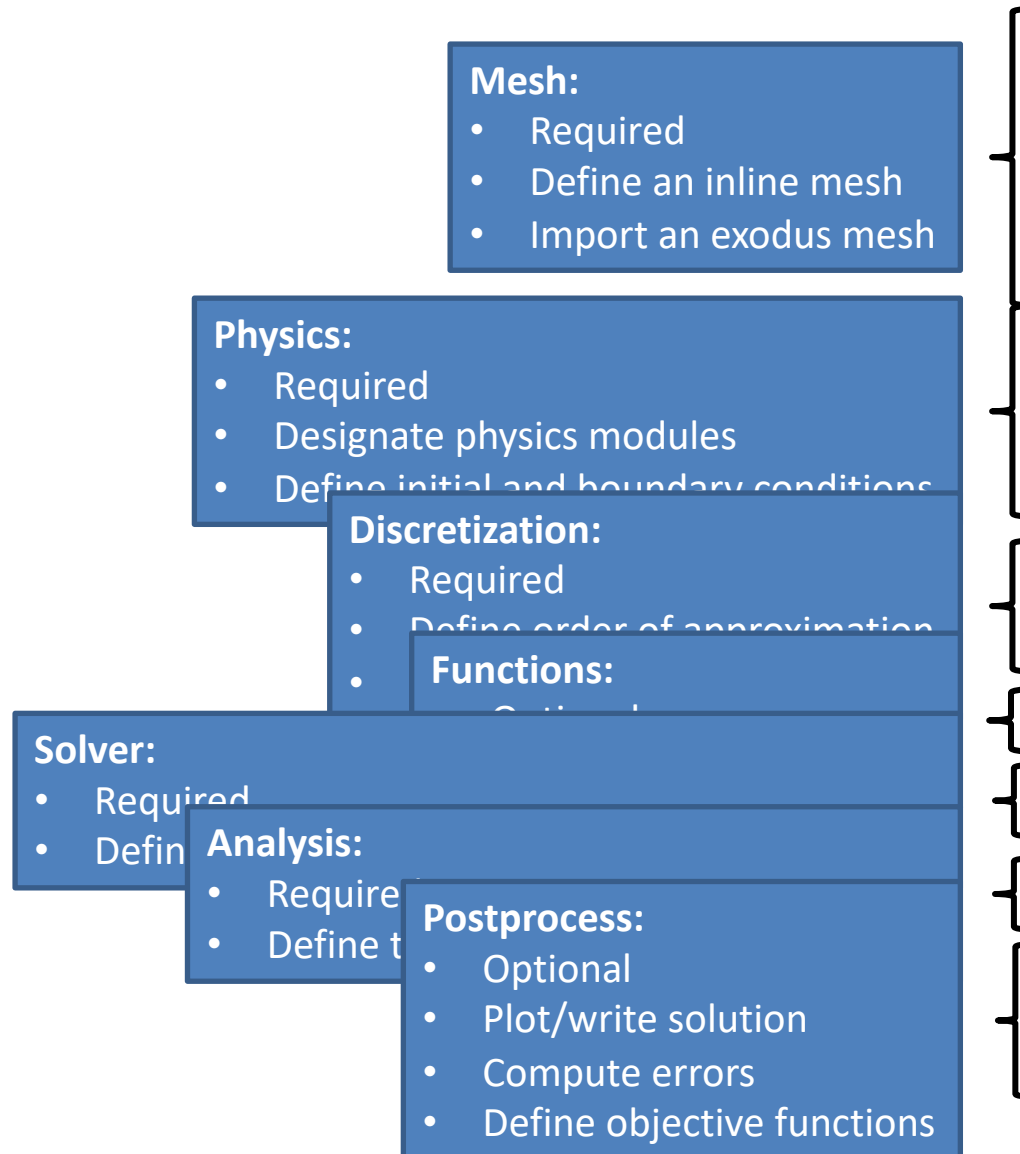
Navigating the User Interface and Input File

- MrHyDE primarily uses the YAML format
 - XML is also an option
- **YAML** (a [recursive acronym](#) for "YAML Ain't Markup Language") is a [human-readable data-serialization language](#).¹
- The user interface looks for input.yaml
- For examples, see the regression tests
- Also, see MrHyDE/scripts/input-files for all the available options.
- There are 5 required blocks and 4 options blocks
- Automatically determines data type
 - int, double, char
- Can force a char by using single quotes

```
YAML 1.1
---
name:
  name:
    dimension: 1
    element type: int
    value: 0.0
    size: 1.0
    price: 0.0
    price: 1.0
    m: 40
    m: 40
  module: thermal
  boundary condition:
    all boundaries: "0.0"
  initial condition:
    m: "0.0"
  discretization:
    order:
      m: 1
    quadrature: 1
  functions:
    thermal module: [calcTemp, calcTemp, calcTemp, calcTemp]
  solver:
    solver: steadystate
    workset size: 100
  analysis:
    analysis type: Forward
  postprocessor:
    compute errors: true
    write solution: false
    true: solution:
      m: calcTemp(calcTemp)
```

¹Source: wikipedia

Navigating the User Interface and Input File



```
1  # Example 1.1
2  #
3  # Geometry
4  #
5  Mesh:
6  #
7  # dimensions: 1
8  # element type: tet
9  # order: 1
10 # mesh: 1
11 # nodes: 1
12 #
13 # Physics
14 #
15 # module: thermal
16 # boundary conditions:
17 #
18 # all boundaries: 1
19 # initial conditions:
20 #
21 #
22 # Discretization
23 #
24 # order:
25 #
26 # quadrature: 1
27 #
28 # Functions
29 #
30 # thermal module: 1
31 #
32 # Solver
33 #
34 # solver: steadystate
35 #
36 # Analysis
37 #
38 # analysis type: Forward
39 #
40 # Postprocess
41 #
42 # compute errors: true
43 # write solution: false
44 #
45 #
46 #
47 #
```

How to Get More Output

- verbosity
 - Default: 0
 - >0 print time step information
 - >1 print nonlinear solver information
 - >8 print linear solver information
 - >=10 print Teuchos timer information (very useful)
 - >20 MueLu preconditioner information
- debug level
 - Default: 0
 - 1: print status when going into manager/interface constructors and some functions that are only called once during setup phase
 - 2: also print status when going into other functions that are called many times during setup or run
 - 3: also print vectors, matrices, some Views



The Function Manager

- One of the most important pieces of MrHyDE for a user to understand
- Similar to Phalanx – builds Directed Acyclic Graphs (DAGs)
- Distinguishing feature: an interpreter that turns strings into DAGs
- Can be thought of as an auto-generator of evaluators (although it doesn't use `PHX::evaluator`)
- To add a function:

```
functionManager->addFunction(name, expression, location)
```

“jumanji” “sin($2.0 * x$)” “ip”

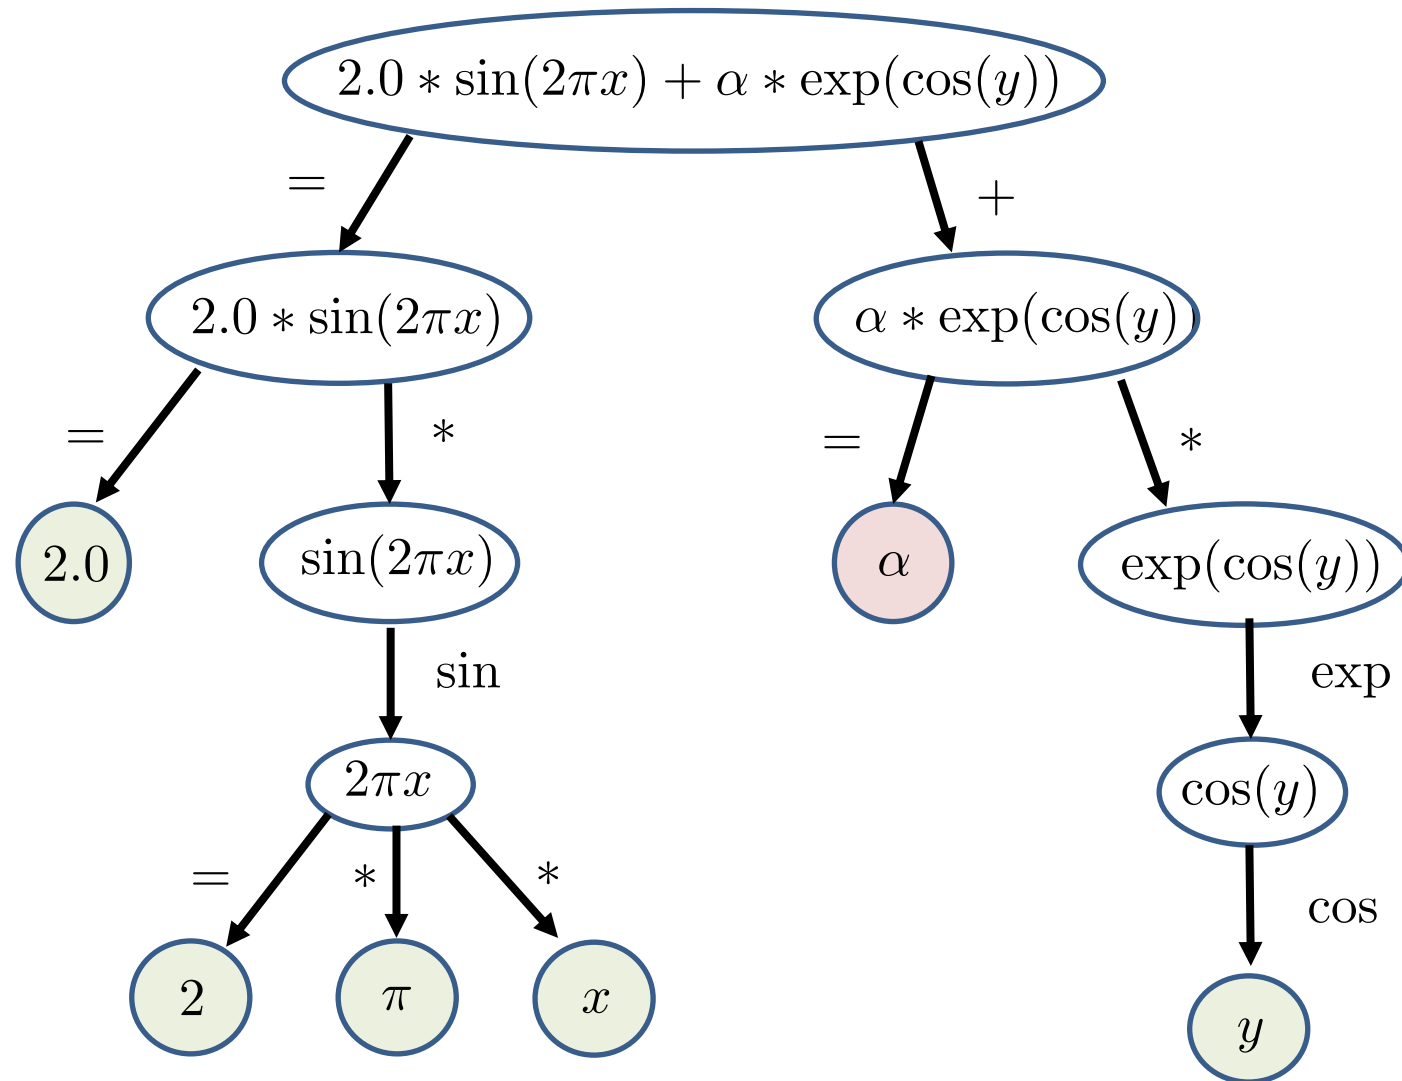
- To evaluate a function:

```
auto data = functionManager->evaluate(name, location)
```

- But, how does it decompose a function into a DAG?

Directed Rooted Trees

Special Case of a Directed Acyclic Graph (DAG)



What can be used in expressions?

Function managers are aware of a few types of variables:

- Spatial variables: x , y , z
- Time: t
- The value of π : π
- Scalars: 2.0, $1.3e-4$, $2.2E-3$, 0.0000001, etc.
- Scalar parameters: λ
- Components of vector parameters: $\eta[1]$
- Discretized parameters: μ
- Solution variables: u , dx , pr , Hu , etc.
- Components of vector variables: $B[x]$
- Components of grad of HGRAD variables: $\text{grad}(dx)[x]$
- Components of curl of HCURL variables: $\text{curl}(E)[y]$
- Divergence of HDIV variables: $\text{div}(B)$
- Names of other functions, responses, objectives: source , obj0
- Normals on sides/faces: n_x , n_y , n_z

Functions Sublist can be Arbitrarily Complicated¹



¹Within reason. There is a limit on the number of recursions, but this can be increased.

Exercise: Create a New Physics Module

- We are going to implement the following PDE:

$$-\Delta \{\text{varname}\} + c \{\text{varname}\} = s(x)$$

- Step 1: choose a name for your variable, e.g., “llama” or “coconut”
- Step 2: choose a name for the physics module, e.g., “coconuts”
- Step 3: copy the template (newmodule.hpp) from MrHyDE/doc/Tutorial/Example into MrHyDE/src/physics

```
cd MrHyDE/src/physics  
cp ../../doc/Tutorial/Example/newmodule.hpp coconuts.hpp
```

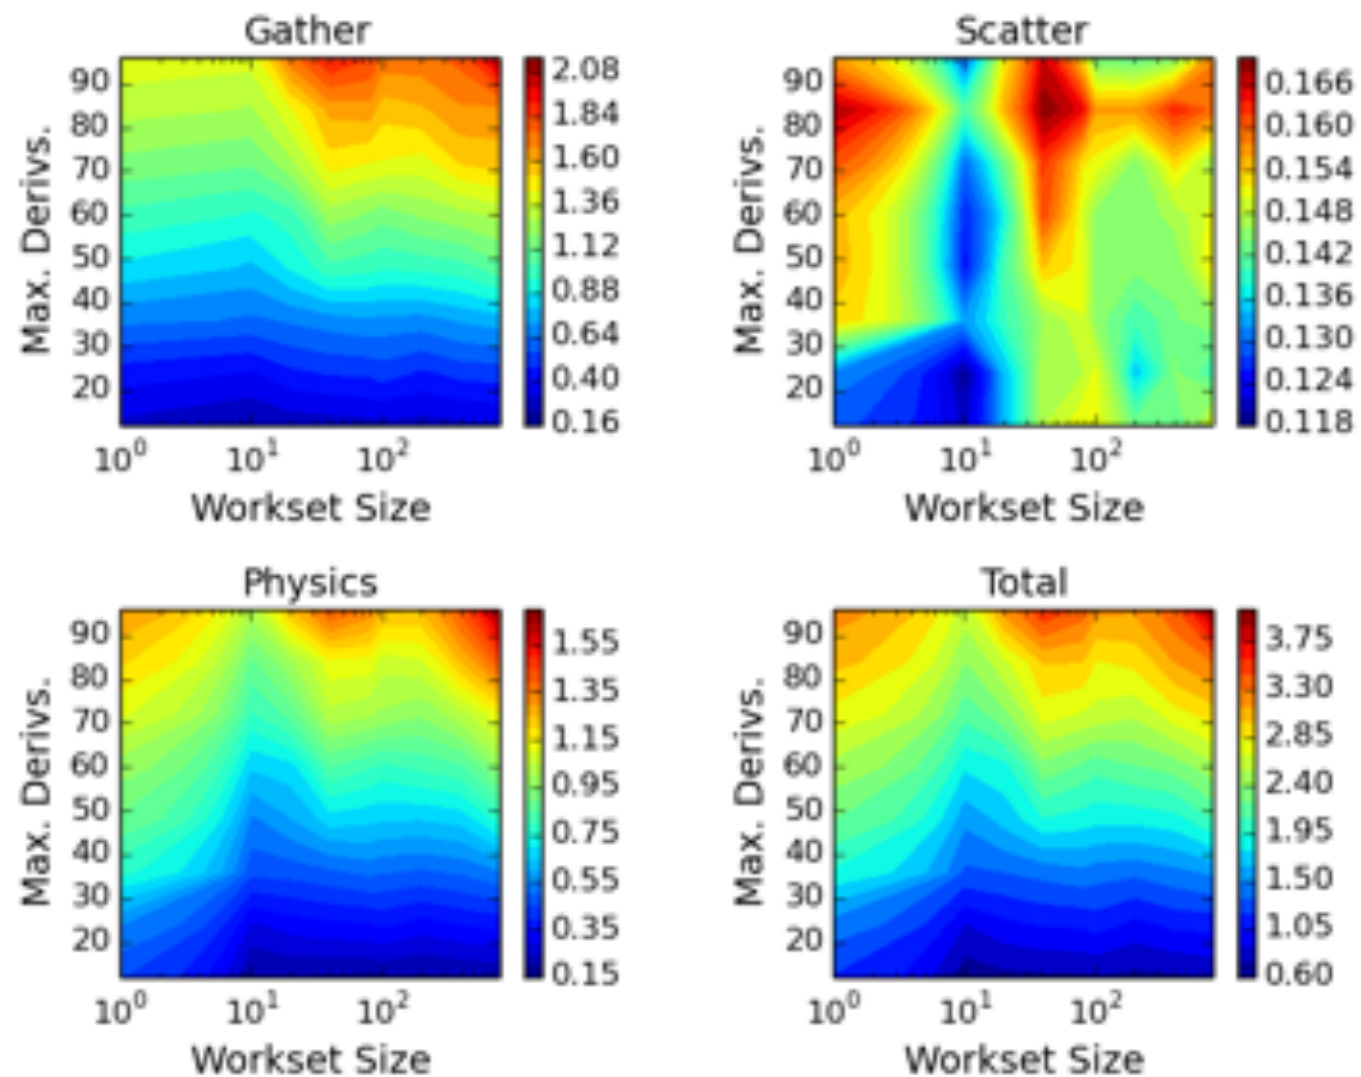
- Step 4: open the file and edit the variable and module names. The weak residual is already there.
- Step 5: make MrHyDE aware of the new module. Open physicsImporter.cpp and add the header and the constructor for your new module (just copy, paste and edit an existing one)
- Step 6: recompile MrHyDE
- Step 7: go back to MrHyDE/doc/Tutorial/Example and edit input.yaml
- Step 8: Create a soft link to MrHyDE/build/src/mrhyde
- Step 9: Run your new module
- Step 10: visualize the results, change the PDE from the input file, couple with other modules, etc.

Boosting Performance

- The default settings are suboptimal to allow all the regression tests to run
- For a specific problem, performance can be optimized without modifying the code
- The default workset size is 100.
 - This is a runtime option to define the number of elements that get processed together
 - 1 is almost never optimal
 - Larger values increase memory requirements
 - Even without threading, adjusting this can improve performance
 - Optimal number is problem dependent
- The default number of derivatives for the SFAD AD objects is 64
 - This is a compile time option (due to Sacado SFAD)
 - Tailoring this can significantly improve performance
 - Minimum value is the maximum of: DOFs per element, active parameters, discretized parameter DOFs per element
 - For example, for the shallow water equations in 2D using linear basis, maxDerivs = 12
 - Adjusted in the MrHyDE configure script:

```
-D MrHyDE_MAX_DERIVS=64 \
```


Impact of Performance Tuning



Preview of Tomorrow

Day 1 - Introduction to Trilinos

- High-level overview of Trilinos
 - *An appropriate build of Trilinos will be available for anyone on the HPC systems. We will not be building Trilinos in this session. If someone does not have access to the HPC systems, I will work with them beforehand to get a build of Trilinos on their Mac or Linux machine.*
- Deeper dive into Kokkos and Sacado.
 - *A basic understanding of these packages will be helpful for day 2.*
- Exercise: creating and working with arrays (Kokkos Views) and automatic differentiation objects (Sacado AD)

Day 2 - Introduction to MrHyDE

- High-level overview of MrHyDE
- How to download, compile, run and visualize results
- Exercise: adding a new PDE in MrHyDE

Day 3 - More advanced features in Trilinos/MrHyDE

- Hybridized methods and concurrent multiscale modeling
- Solving coupled multiphysics problems
- Performance portability and using heterogeneous computational architectures
- Large-scale PDE constrained optimization